

COSC 5557: Practical ML

Warm Up Exercise

Abiodun Awosola

2024-03-02

Loading Basic Packages

```
knitr::opts_chunk$set(comment = NA) # removes '##' from outputs

# .Rprofile  (# allows universal package installation)
options(repos = c(
  mlrorg = "https://mlr-org.r-universe.dev",
  CRAN = "https://cloud.r-project.org/"
))

library(mlr3verse)
library(mlr3learners)

lgr::get_logger("mlr3")$set_threshold("warn")
```

Creating Tasks and Learners

```
# Imports Data

wine_data <- read.table("winequality-white.csv", sep = ";",
  check.names = TRUE, header=T)

str(wine_data)

'data.frame':  4898 obs. of  12 variables:
 $ fixed.acidity      : num  7 6.3 8.1 7.2 7.2 8.1 6.2 7 6.3 8.1 ...
 $ volatile.acidity   : num  0.27 0.3 0.28 0.23 0.23 0.28 0.32 0.27 0.3 0.22 ...
 $ citric.acid        : num  0.36 0.34 0.4 0.32 0.32 0.4 0.16 0.36 0.34 0.43 ...
 $ residual.sugar     : num  20.7 1.6 6.9 8.5 8.5 6.9 7 20.7 1.6 1.5 ...
 $ chlorides          : num  0.045 0.049 0.05 0.058 0.058 0.05 0.045 0.045 0.049 0.044 ...
 $ free.sulfur.dioxide : num  45 14 30 47 47 30 30 45 14 28 ...
 $ total.sulfur.dioxide: num  170 132 97 186 186 97 136 170 132 129 ...
 $ density            : num  1.001 0.994 0.995 0.996 0.996 ...
 $ pH                 : num  3 3.3 3.26 3.19 3.19 3.26 3.18 3 3.3 3.22 ...
 $ sulphates          : num  0.45 0.49 0.44 0.4 0.4 0.44 0.47 0.45 0.49 0.45 ...
 $ alcohol            : num  8.8 9.5 10.1 9.9 9.9 10.1 9.6 8.8 9.5 11 ...
 $ quality            : int   6 6 6 6 6 6 6 6 6 6 ...
```

```

# Creates mlr3 task;
# target is the column to be learnt

wine_tsk = as_task_classif(wine_data, target = "quality")

print(wine_tsk)

<TaskClassif:wine_data> (4898 x 12)
* Target: quality
* Properties: multiclass
* Features (11):
  - dbl (11): alcohol, chlorides, citric.acid, density, fixed.acidity,
    free.sulfur.dioxide, pH, residual.sugar, sulphates,
    total.sulfur.dioxide, volatile.acidity

```

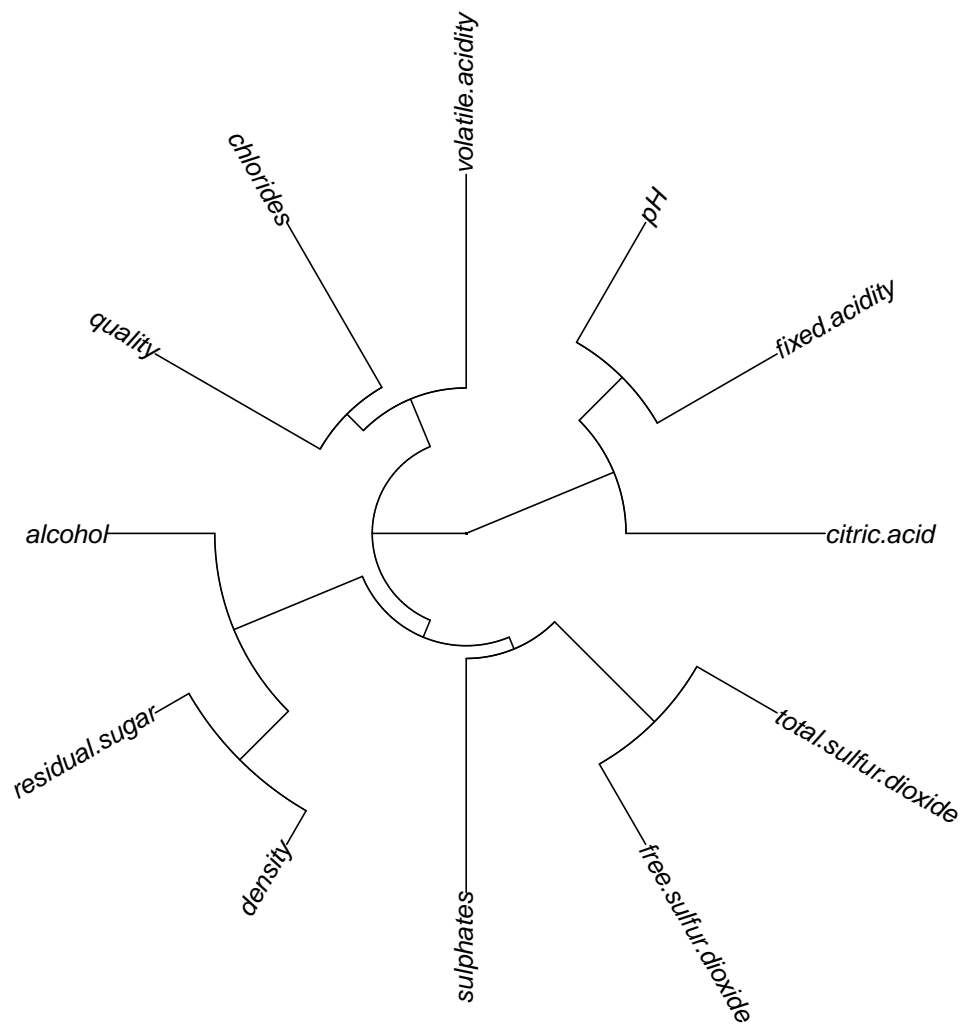
Phylogenetic tree

```

# phylogenetic tree

library(ClustOfVar)
library(ape)
vctree <- hclustvar(wine_data)
plot(as.phylo(vctree), type = "fan")

```



```
# creates learner  
# 2 equivalent calls:  
  
learner_1 = mlr_learners$get("classif.rpart")  
learner_1 = lrn("classif.rpart")  
print(learner_1)
```

```
<LearnerClassifRpart:classif.rpart>: Classification Tree  
* Model: -
```

```

* Parameters: xval=0
* Packages: mlr3, rpart
* Predict Types: [response], prob
* Feature Types: logical, integer, numeric, factor, ordered
* Properties: importance, missings, multiclass, selected_features,
  twoclass, weights

```

Train and Predict

```

# trains learner on subset of task
learner_1$train(wine_tsk, row_ids = 1:3918)

# this is what the decision tree looks like
print(learner_1$model)

```

```
n= 3918
```

```

node), split, n, loss, yval, (yprob)
  * denotes terminal node

```

```

1) root 3918 2237 6 (0.0048 0.035 0.3 0.43 0.19 0.04 0.0013)
 2) alcohol< 10.85 2560 1438 6 (0.0043 0.044 0.41 0.44 0.091 0.011 0.00039)
   4) volatile.acidity>=0.2425 1466 703 5 (0.0048 0.063 0.52 0.36 0.046 0.0041 0.00068)
     8) alcohol< 9.75 897 353 5 (0.0045 0.057 0.61 0.3 0.03 0.0011 0) *
     9) alcohol>=9.75 569 310 6 (0.0053 0.072 0.38 0.46 0.072 0.0088 0.0018)
       18) free.sulfur.dioxide< 17.5 108 53 5 (0.0093 0.2 0.51 0.23 0.046 0 0) *
       19) free.sulfur.dioxide>=17.5 461 227 6 (0.0043 0.041 0.36 0.51 0.078 0.011 0.0022) *
   5) volatile.acidity< 0.2425 1094 501 6 (0.0037 0.018 0.27 0.54 0.15 0.02 0) *
 3) alcohol>=10.85 1358 799 6 (0.0059 0.019 0.1 0.41 0.37 0.094 0.0029)
   6) alcohol< 12.55 1100 612 6 (0.0064 0.02 0.12 0.44 0.33 0.081 0.0018) *
   7) alcohol>=12.55 258 124 7 (0.0039 0.016 0.031 0.28 0.52 0.15 0.0078) *

```

```

# predicts using observations from task
prediction = learner_1$predict(wine_tsk, row_ids = 3919:4898)
print(prediction)

```

```
<PredictionClassif> for 980 observations:
```

row_ids	truth	response
3919	7	7
3920	6	7
3921	6	6

4896	6	6
4897	7	7
4898	6	6

Evaluation

Scoring the Prediction object with some metrics. And take a deeper look by inspecting the confusion matrix.

```
head(as.data.table(mlr_measures))
```

	key	label	task_type	packages
1:	aic	Akaike Information Criterion	<NA>	mlr3

```

2:          bic Bayesian Information Criterion      <NA>          mlr3
3:   classif.acc      Classification Accuracy   classif mlr3,mlr3measures
4:   classif.auc      Area Under the ROC Curve   classif mlr3,mlr3measures
5:   classif.bacc      Balanced Accuracy         classif mlr3,mlr3measures
6: classif.bbrier      Binary Brier Score        classif mlr3,mlr3measures

```

```

predict_type task_properties

```

```

1:      <NA>

```

```

2:      <NA>

```

```

3:   response

```

```

4:     prob      twoclass

```

```

5:   response

```

```

6:     prob      twoclass

```

```

scores = prediction$score(msr("classif.acc"))

```

```

print(scores)

```

```

classif.acc

```

```

0.5877551

```

```

scores = prediction$score(msrs(c("classif.acc", "classif.ce")))

```

```

print(scores)

```

```

classif.acc  classif.ce

```

```

0.5877551    0.4122449

```

Confusion matrix

```

cm = prediction$confusion

```

```

print(cm)

```

```

      truth
response 3  4  5  6  7  8  9
3      0  0  0  0  0  0  0
4      0  0  0  0  0  0  0
5      0 10 121 55  0  0  0
6      1 15 144 414 110 13  0
7      0  0  1  48  41  7  0
8      0  0  0  0  0  0  0
9      0  0  0  0  0  0  0

```

Key to understand the confusion matrix

- 5 was predicted as 6 144 times. $(\text{response}, \text{truth}) = (6, 5)$

Changing Hyperparameters

The **Learner** contains information about all parameters that can be configured, including data type, constraints, defaults, etc. The hyperparameters can be changed either during construction or later through an active binding.

```

as.data.table(learner_1$param_set)[, .(id, class, lower, upper, nlevels)]

```

```

      id    class lower upper nlevels
1:      cp ParamDbl    0     1      Inf
2: keep_model ParamLgl  NA    NA       2
3:  maxcompete ParamInt    0   Inf      Inf

```

```

4:      maxdepth ParamInt    1    30    30
5:  maxsurrogate ParamInt    0   Inf   Inf
6:      minbucket ParamInt    1   Inf   Inf
7:      minsplit ParamInt    1   Inf   Inf
8: surrogatestyle ParamInt    0    1    2
9:   usesurrogate ParamInt    0    2    3
10:          xval ParamInt    0   Inf   Inf

learner_2 = lrn("classif.rpart", predict_type = "prob", minsplit = 50)
learner_2$param_set$values$minsplit = 50

```

Resampling

Resampling repeats the train-predict-score loop and collects all results in a nice `'data.table::data.table()'`.

```

cv10 = rsmp("cv", folds = 10)
rr = resample(wine_tsk, learner_1, cv10)
print(rr)

```

```

<ResampleResult> with 10 resampling iterations
  task_id learner_id resampling_id iteration warnings errors
wine_data classif.rpart          cv          1          0          0
wine_data classif.rpart          cv          2          0          0
wine_data classif.rpart          cv          3          0          0
wine_data classif.rpart          cv          4          0          0
wine_data classif.rpart          cv          5          0          0
wine_data classif.rpart          cv          6          0          0
wine_data classif.rpart          cv          7          0          0
wine_data classif.rpart          cv          8          0          0
wine_data classif.rpart          cv          9          0          0
wine_data classif.rpart          cv         10          0          0

```

```
rr$score(msr(c("classif.acc", "classif.ce"))[, .(iteration, task_id, learner_id, resampling_id, classif.ce)])
```

```

  iteration task_id learner_id resampling_id classif.ce
1:         1 wine_data classif.rpart          cv 0.4714286
2:         2 wine_data classif.rpart          cv 0.5000000
3:         3 wine_data classif.rpart          cv 0.4795918
4:         4 wine_data classif.rpart          cv 0.4387755
5:         5 wine_data classif.rpart          cv 0.5204082
6:         6 wine_data classif.rpart          cv 0.4836735
7:         7 wine_data classif.rpart          cv 0.4061224
8:         8 wine_data classif.rpart          cv 0.4571429
9:         9 wine_data classif.rpart          cv 0.4580777
10:        10 wine_data classif.rpart          cv 0.4723926

```

```

# gets all predictions nicely concatenated in a table
prediction = rr$prediction()
as.data.table(prediction)

```

```

  row_ids truth response
1:      2      6        5
2:      6      6        6
3:     17      6        5
4:     43      6        5
5:     52      7        6

```

```

---
4894:    4838      6      7
4895:    4845      6      6
4896:    4848      7      6
4897:    4865      5      5
4898:    4898      6      6

# The confusion matrix for entire prediction
cm = prediction$confusion
print(cm)

```

```

      truth
response 3  4  5  6  7  8  9
      3  0  0  0  0  0  0
      4  0  0  0  0  0  0
      5  5 85 787 439 41  1  0
      6 14 74 661 1640 664 129 3
      7  1  4  9 119 175 45  2
      8  0  0  0  0  0  0  0
      9  0  0  0  0  0  0  0

```

Populating the learner dictionary

mlr3learners ships out with a dozen different popular Learners. They can be listed from the dictionary. If more were desired, an extension package, mlr3extralearners, could be installed from GitHub. Importantly, after loading mlr3extralearners, the dictionary increases in size.

```
head(as.data.table(mlr_learners)[, c("key", "packages")])
```

	key	packages
1:	classif.AdaBoostM1	mlr3,mlr3extralearners,RWeka
2:	classif.C50	mlr3,mlr3extralearners,C50
3:	classif.IBk	mlr3,mlr3extralearners,RWeka
4:	classif.J48	mlr3,mlr3extralearners,RWeka
5:	classif.JRip	mlr3,mlr3extralearners,RWeka
6:	classif.LMT	mlr3,mlr3extralearners,RWeka

```
library(mlr3extralearners)
print(as.data.table(mlr_learners)[, c("key", "packages")])
```

	key	packages
1:	classif.AdaBoostM1	mlr3,mlr3extralearners,RWeka
2:	classif.C50	mlr3,mlr3extralearners,C50
3:	classif.IBk	mlr3,mlr3extralearners,RWeka
4:	classif.J48	mlr3,mlr3extralearners,RWeka
5:	classif.JRip	mlr3,mlr3extralearners,RWeka

172:	surv.ranger	mlr3,mlr3proba,mlr3extralearners,ranger
173:	surv.rfsrc	mlr3,mlr3proba,mlr3extralearners,randomForestSRC,pracma
174:	surv.rpart	mlr3,mlr3proba,rpart,distr6,survival
175:	surv.svm	mlr3,mlr3proba,mlr3extralearners,survivalsvm
176:	surv.xgboost	mlr3,mlr3proba,mlr3extralearners,xgboost

Benchmarking multiple learners

The `benchmark` function can conveniently compare ‘r ref(“Learner”, “Learners”) on the same dataset(s).

```
learners = list(learner_1, learner_2, lrn("classif.randomForest"))
grid = benchmark_grid(wine_tsk, learners, cv10)
bmr = benchmark(grid)
print(bmr)
```

<BenchmarkResult> of 30 rows with 3 resampling runs

nr	task_id	learner_id	resampling_id	iters	warnings	errors
1	wine_data	classif.rpart	cv	10	0	0
2	wine_data	classif.rpart	cv	10	0	0
3	wine_data	classif.randomForest	cv	10	0	0

```
print(bmr$aggregate(measures = msrs(c("classif.acc", "classif.ce"))))
```

	nr	task_id	learner_id	resampling_id	iters	classif.acc	classif.ce
1:	1	wine_data	classif.rpart	cv	10	0.5320521	0.4679479
2:	2	wine_data	classif.rpart	cv	10	0.5320521	0.4679479
3:	3	wine_data	classif.randomForest	cv	10	0.7051863	0.2948137

Hidden columns: resample_result